Complete image actor: execution verified

7.548B total parameters 25.58M trainable 17.181 GiB peak allocated

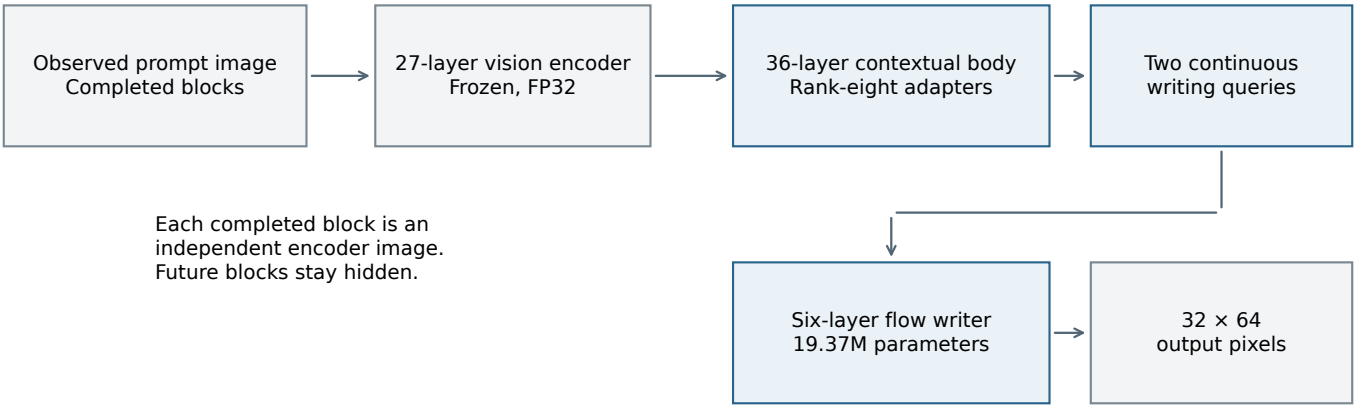

$$Y_t = (1-t)U + tY, \quad U \sim \mathcal{N}(0, I), \quad \mathcal{L}_{\text{flow}} = \mathbb{E}\|v_\theta(Y_t, t, h) - (Y - U)\|^2$$

Frozen contextual weights use BF16 storage and FP32 computation. A separate head predicts stopping.

No language vocabulary or external model call in the actor. Writing and language quality remain unqualified.