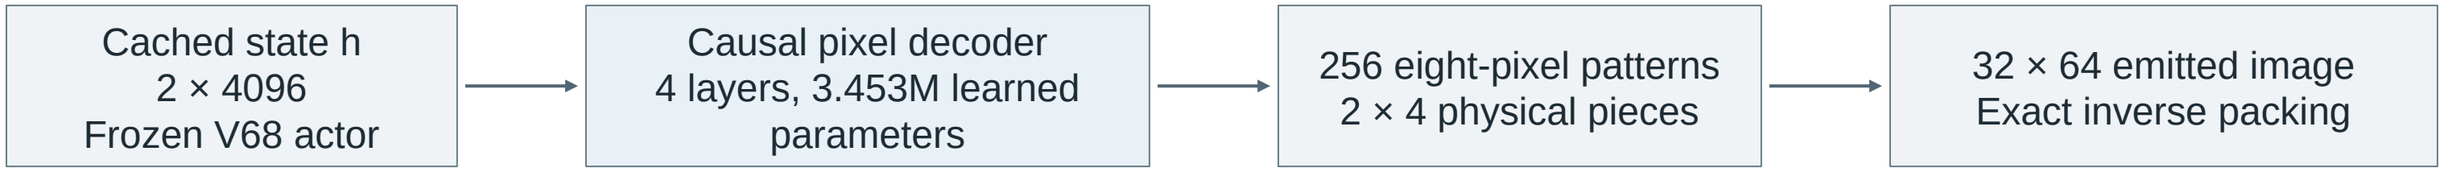
Direct pixel likelihood fits known blocks but fails held-out writing

a Conditional physical-image density

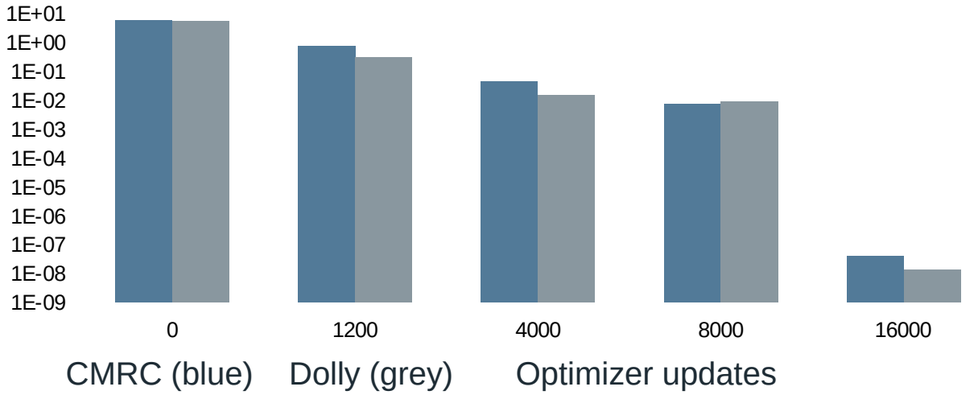


Training supplies preceding target pieces. Generation uses only its own emitted pieces. No lexical output inventory or font renderer.

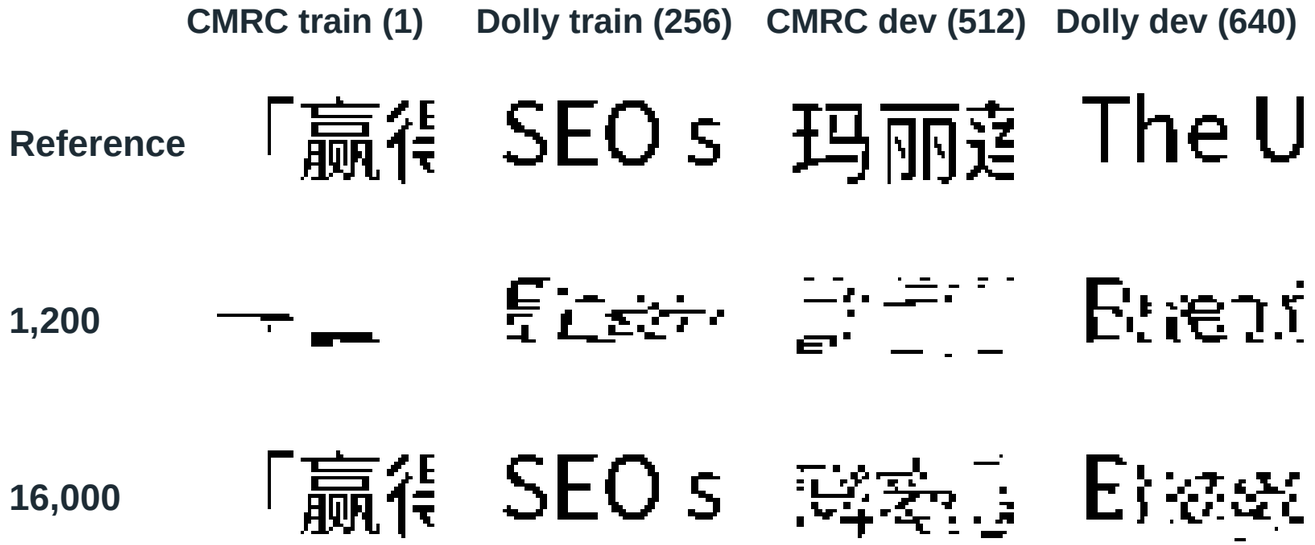
b Exact autonomous first blocks

Cohort / updates	1,200	16,000
CMRC training	0/256	256/256
Dolly training	12/256	256/256
CMRC development	0/128	0/128
Dolly development	0/64	0/64

Whole-training-set NLL per eight-pixel piece (log scale)



c Reference, shorter fit and longer fit



Original 32 × 64 physical blocks. Boundary-clipped glyphs are retained.

Longer fit: 4.65 min, 0.825 GiB head peak. Costs exclude the pretrained reader and cache. All 896 outputs scored; training memorization does not establish language generalization.