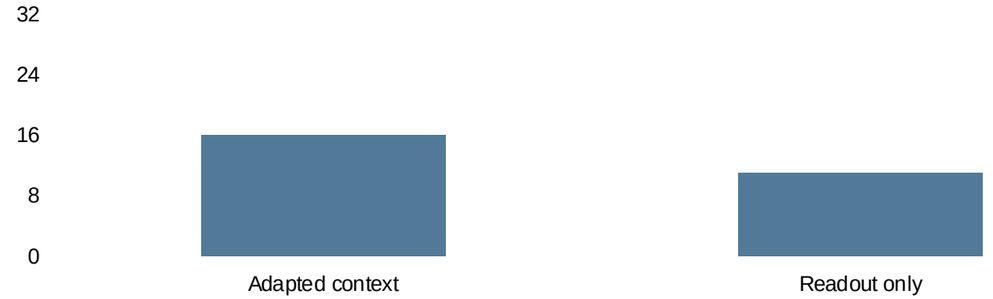
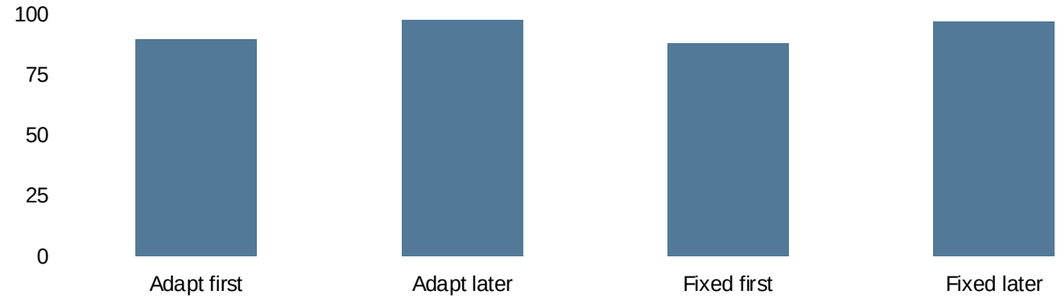
Contextual adaptation improves image answers; reliability remains limited

a Clear complete answers on new contexts



Both questions correct: 5/16 vs 2/16 contexts.
Neither meets the fixed 24/32 and 10/16 pilot.

b First/later training foreground F1 (%)



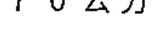
First: 3,176; later: 20,086 reference-conditioned targets.
The first image is decided before any generated feedback.

c Image prompt → **V + learned context** → **Continuous G + image D** → **Actual pixels -> E -> next state**

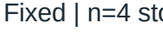
d First eight reserved questions, both conditions, all outcomes retained

Question 00

Adapted | n=4 stop

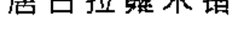


Fixed | n=4 stop

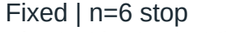


Question 02

Adapted | n=6 stop

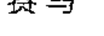


Fixed | n=6 stop

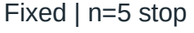


Question 04

Adapted | n=2 stop

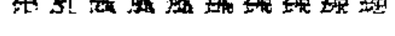


Fixed | n=5 stop



Question 06

Adapted | n=64 cap | first 10

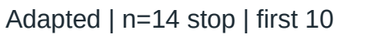


Fixed | n=64 cap | first 10

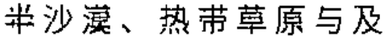


Question 01

Adapted | n=14 stop | first 10

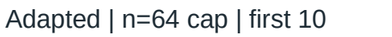


Fixed | n=64 cap | first 10

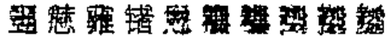


Question 03

Adapted | n=64 cap | first 10



Fixed | n=64 cap | first 10



Question 05

Adapted | n=9 stop



Fixed | n=9 stop

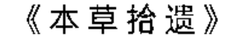


Question 07

Adapted | n=6 stop



Fixed | n=6 stop



Unrepaired output pixels. Prefixes only where marked; full PNGs embedded. Agent reading is not independent human grading. Answer pilot remains failed.