Learning conditional language in continuous image space

Imagized Language Model Project

Writing includes historical and handwritten forms that cannot always be specified by a character code. We investigate conditional language generation directly between writing images on one consumer GPU. Runtime inference uses physical rasters without a character vocabulary. A frozen continuous codec reconstructs 126 of 128 development blocks exactly, at 99.998% foreground F1. A fresh 64.520M-parameter writer trained on 12,976 Chinese/English examples yields only two exact Chinese answers out of 128 and fails all 12 new copying requests. Retaining full native context also fails: training 46.900M interface parameters over 3,193 Chinese examples for 40,000 updates gives no exact development answers, at 36.504% F1. Erasing the question reduces F1 to 25.404%, showing input sensitivity without accurate answers. On 32 training questions, correct answer history and length still give no exact images at 50.241% F1, while supplied-code decoding reaches 99.9953%. Conditional prediction therefore fails before generated-history errors can explain the outcome. The complete native actor has 7.583B parameters and retains externally pretrained weights; local costs exclude that pretraining. Decoder-sensitive loss also fails, lowering free training-image F1 from 85.984% to 70.009%, with no exact answers. These controls separate rendering, contextual information and autonomous writing. Useful image-language modelling remains unachieved.

Written language is encountered as marks on paper, screens and historical surfaces. Its visual form carries information about script, handwriting and material history. Modern character encodings provide a powerful interface to many such marks, but a finite inventory cannot express every damaged, transitional or unencoded form. An image-based language model could retain that evidence throughout reading and generation, with transcription offered as an optional output view.

Pixel-based language modelling already has substantial precedents. PIXAR demonstrates autoregressive language generation in pixel space and identifies the difficulty of generating readable text[1]. Continuous autoregression can use a diffusion loss in place of vector quantization[2], while Diffusion Forcing models continuous sequences under independently chosen noise levels[3]. These results motivate a visual language interface, but do not establish that a small glyph generator acquires linguistic knowledge or that image input automatically reduces training cost. Exact pixel factorization also has established precedents in PixelRNN and conditional PixelCNN[6,7]. Masked autoregressive networks provide normalized probabilities without a character dictionary[8]. Multilingual generation from pixels has also been studied[9].

Recent continuous multimodal flow models share a generative objective across embeddings while retaining token-derived language encoders and text decoders[15]. OCG reports rendered-question and answer-image reasoning with teacher assistance on the restricted CLEVR task[16]. We focus on independent image-only inference and controlled tests linking conditional scores to autonomous writing.

We investigate the conditional generation problem directly. The model receives raster history, predicts a distribution of future writing, emits pixels and rereads those pixels. Chinese provides a demanding initial case: stroke configurations must remain coherent while context selects among many similar written forms. The present study uses one canonical font to isolate this coupling before introducing handwriting or historical scripts.

Results

A continuous visual interface preserves the written signal

Each observed cell is a 32×32 ink image. An orthonormal transform maps its signed pixels into a 1,024-dimensional continuous field. The transform is invertible and contains no learned character dictionary. An ordered sequence of these fields enters a causal transformer. The deployed model receives no Unicode values, OCR transcripts, character IDs or vocabulary logits (Fig. 1). Host-side rendering uses encoded source text during data preparation. This distinction matters: the current benchmark studies learning from rendered Chinese, rather than discovering a script from raw archaeological photographs.

Ordered image history contains next-writing information

The strongest established reader result uses 2,048 deterministic development windows and an evaluator-only bank of 1,024 canonical glyph images. Full-history next-image top-1 is 19.97%, above the image unigram at 1.42%, the symbolic bigram at 12.26% and the prefix-shuffled condition at 18.36%. The ordered history improves target log probability by 0.174 nat over the shuffled control. This result supports bounded conditional prediction from image sequences.

These figures measure selection among completed images. The evaluator bank does not enter the model and cannot act as a deployed writer. Thresholding the predicted continuous field can mix strokes from several plausible futures. For squared-error prediction, the population optimum is $f^*(H) = \mathbb{E}[X \mid H]$. A conditional mean need not be a valid glyph when several futures are possible. Sharp copying, target ranking and sampled language generation therefore require separate measurements.

A glyph prior does not automatically bind to language

To separate visual synthesis from contextual prediction, we train an unconditional glyph prior and a noisy visual compatibility critic while freezing the causal reader. The combined runtime contains 23.27 million parameters. The critic scores an arbitrary noisy image against the visible history. In the ideal density-ratio limit,

$$\nabla_x \log p_t(x \mid H) = \nabla_x \log p_t(x) + \nabla_x \log \frac{p_t(x \mid H)}{p_t(x)}. \qquad (1)$$

The construction is related to classifier-guided diffusion[4].

On 512 development windows, the clean critic reaches 21.875% top-1, versus 20.898% with shuffled history. Its first guided sample reaches zero target identities in 256 cases. Select-

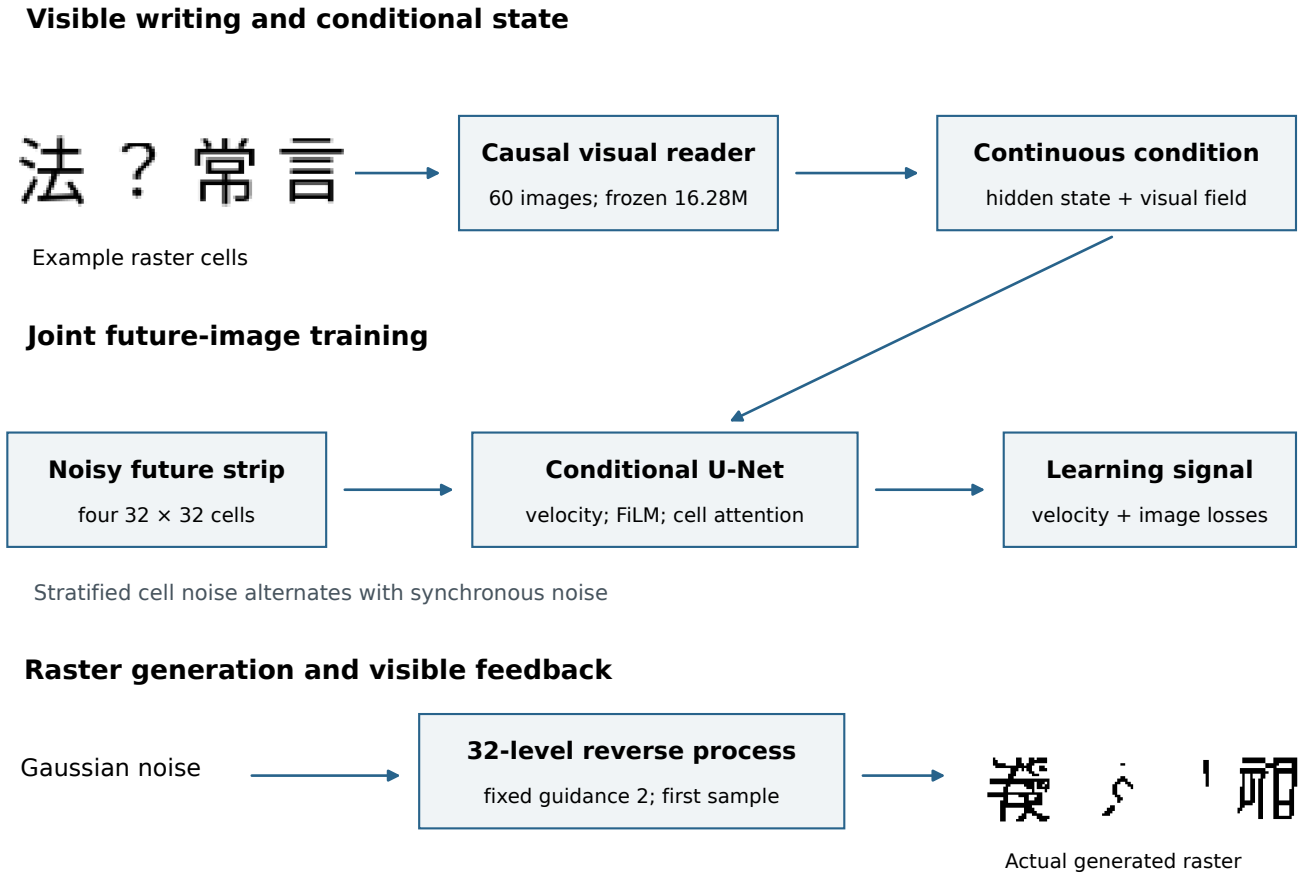


Fig. 1 | A visual language interface and its conditional generation test. a, Rendered or native writing supplies ordered image cells to a causal visual reader. Text rendering is an optional host interface. **b**, A conditional denoiser uses the frozen reader state and a noisy four-cell strip to predict diffusion velocity. Future cells exchange information through convolutions and pooled-cell attention. **c**, Sampling starts from noise, emits four visible cells and appends their pixels to the next context. Optional transcription follows generation and is absent from the model. The schematic shows the implemented experimental contract, without claiming that its generated content is already correct.

ing the highest clean critic score among four samples reaches 1/256, with the same result after shuffling history. The generated glyph-like marks are diverse and nonblank, yet fail conditional identity (Fig. 2).

This experiment identifies a reader–writer gap without proving a unique cause. At the terminal corruption level, target signal power is 2.43×10^{-7}, and an exact conditional guidance term should vanish. That limit also occurs in successful conditional diffusion. It cannot alone explain the failure. Insufficient conditional-score accuracy, off-manifold gradients, limited prior quality and finite sampling remain compatible explanations. We consequently test a directly conditional writer rather than interpret clean-image ranking as evidence of a generative solution.

Direct conditional block generation tests the missing coupling

The new writer models a joint four-cell future conditioned on 60 visible cells. It reuses the 16.28M-parameter reader and fine-tunes a 5.90M-parameter strip denoiser. A projected reader state modulates each residual scale, and a four-cell attention layer exchanges future information at the bottleneck. There is no separate guidance critic or candidate selector.

For signed clean writing x_0, noise ε and cumulative signal $\bar{\alpha}_t$, the training target is velocity:

$$x_t = \sqrt{\bar{\alpha}_t}\,x_0 + \sqrt{1-\bar{\alpha}_t}\,\varepsilon, \tag{2}$$

$$v^* = \sqrt{\bar{\alpha}_t}\,\varepsilon - \sqrt{1-\bar{\alpha}_t}\,x_0. \tag{3}$$

The clean estimate is $\hat{x}_0 = \sqrt{\bar{\alpha}_t}\,x_t - \sqrt{1-\bar{\alpha}_t}\,\hat{v}$. Training alternates independently stratified future noise levels with synchronous levels that cover the sampling distribution. Fixed visual-field and foreground losses retain the writing signal. Intermediate target-field alignment follows the motivation of representation alignment[5], using a fixed raster transform rather than a semantic teacher.

The run completes all 8,000 updates in 361.6 s, using 0.543 GiB peak allocated CUDA memory. Evaluation uses one primary sample from 256 prefixes under paired initial noise. Ordered first-position identity is 1.17%, compared with 1.17% after shuffling and 0.00% with a zero condition. Across all four future cells, identity is 0.586%, with an ordered-minus-shuffled difference of 0.098 percentage points and a record-bootstrap 95% interval of [0.000, 0.307]. Pixel F1 is 0.320. The registered complete-mechanism decision is `reject`. The raw image panel shows the first fixed examples without reranking or selection (Fig. 2).

A post-result motor diagnostic supplies the intended first-cell image to the same trained writer on 64 fixed prefixes. Exact plan and image-field supply raise pixel F1 from 0.378 to 0.459 and target identity from 1/64 to 5/64. The fixed inverse transform

reproduces the supplied raster with F1 1.0. Because this test supplies the answer, it provides no autonomous-language evidence. It shows that the current learned writer still loses substantial shape information even with an exact visual condition.

Exact raster likelihood preserves copying but not conditional language

The diffusion writer's supplied-image failure motivates a distribution that retains the exact inverse visual field as aligned pixel logits. For binary image X, ordered history H and row-major pixel index p, we define

$$p_\theta(X \mid H) = \prod_{p=1}^{1024} \mathrm{Bern}[X_p; \sigma(6A_p + \delta_\theta(H, X_{<p}))]. \qquad (4)$$

Here A is the inverse predicted field, scaled to signed pixels and clipped to $[-2,2]$. A causal row transformer and within-row MADE supply the residual. The density contains 2,434,592 learned parameters, combined with the frozen reader for 18,711,841 runtime parameters. This is a conditional raw-pixel factorization [6–8], not a new language vocabulary.

With a supplied binary image, A_p is exactly ± 1. A zero-initialized residual therefore copies its greedy mode exactly before training. This is an architectural property, not acquired linguistic knowledge. Pure image NLL training retains that property: greedy copying is exact for all 256 development targets, stochastic-copy F1 is 0.99949, and a deranged supplied image reduces F1 to 0.358 (Fig. 3). These exercises supply the answer and cannot qualify autonomous generation.

On 512 natural targets, ordered NLL is 15.219 nats per image (0.02144 bits per pixel), versus 15.254 with earlier history shuffled. The paired advantage is 0.0355 nat, with source-record bootstrap 95% interval [-0.1014, 0.1693]. This interval includes zero. An independent-pixel distribution using the same predicted plan has NLL 846.09 nats. That comparison supports learned dependencies within the image; it does not isolate learned linguistic dependencies between images.

Autonomous sampling recovers 4/256 nearest target identities, versus 3/256 after shuffling, with pixel F1 0.350. The paired identity advantage is 0.391 percentage points, with interval [0, 1.210]. Greedy generation reaches 14/256 exact images but only 0.215 F1 and 0.330 target-relative ink, consistent with frequent sparse forms in the visible gallery. Four-step feedback identity is 4.297% ordered and 3.125% shuffled; this small secondary difference does not reverse the failed primary language tests. The registered decision is `reject`.

Joint learning improves the conditional image distribution

A matched experiment tests whether the same raster law benefits from learning its reader under image NLL. Both arms start from the exact-density EMA and receive identical new training draws. One continues to freeze the reader; the other differentiates likelihood through its hidden state and predicted field. Reader dropout remains disabled in both, so gradient flow is the controlled change. The unused contrastive temperature stays fixed.

After 8,000 updates per arm, joint target NLL is 7.026 versus 7.579 for the frozen control. The paired gain is 0.553 nat, with interval [0.450, 0.665]. Joint ordered history improves NLL over earlier-shuffled history by 0.185 nat, with interval [0.128, 0.244]. The frozen continuation also gains an earlier-history effect (0.119 nat, interval [0.052, 0.186]); extra training alone therefore explains part of the progress. The matched comparison isolates the additional likelihood benefit of updating the reader.

Joint first-sample identity is 9/256, versus 6/256 for the matched control and 8/256 after shuffling. The joint-minus-control identity difference is 1.172 percentage points, with interval [-0.397, 2.893]. Both arms retain exact greedy copying of all 256 supplied images. Joint sampling has F1 0.362 and isolated-ink fraction 0.00744. The joint arm passes six of eight gates, but fails sampled quality and ordered-generation criteria; its decision remains `reject` (Supplementary Note 13 and Fig. S2). This is evidence of improved conditional learning, not a qualified language generator.

Raw-pixel search improves choice but exposes repetitive rollouts

Known-form validity and target correctness separate two remaining problems. A post-result audit of the original saved images finds exact known-form membership in 43.750% of joint first samples, versus 33.594% for matched frozen training. The paired gain is 10.156 percentage points [4.219, 16.240]. Joint greedy images reach 98.438% known-form membership, but only 7.422% exact target match (Supplementary Note 14). Clear writing alone is insufficient.

We next search the same normalized pixel law with eight partial-raster beams, without a glyph inventory or weight change. On 2,048 new deterministic development windows, ordered exact next-image match is 11.621%, versus 8.252% for original pixel greediness and 9.863% after earlier-history shuffling. The beam-minus-greedy gain is 3.369 points [2.351, 4.412]; the earlier-history gain is 1.758 points [0.966, 2.552]. Every primary beam output matches a known canonical raster after generation.

Eight-step feedback on 128 prefixes reaches 6.641% exact target images, versus 4.883% greedy, but adjacent repeated images rise to 25.446%. The actual target rate is 0.112%. Thus the registered decoder decision is `reject_decoder`: five of six gates pass, but repetitive rollout fails (Supplementary Note 15 and Fig. S4). This outcome is compatible with known limits of probability-maximizing language decoding [10]. No repetition filter or output replacement is applied.

Natural questions expose the gap between selection and complete answers

To test use of language beyond continuation, we add a copy distribution over images in the current prompt, following the established pointer–generator idea [11]. Repeated identical images contribute summed attention mass. A learned gate combines copying with the exact raster law, and a separate stop hazard controls answer length. Only observed pixels and causal states enter this 19.01M-parameter model. No reference answer, symbolic identity or persistent writing inventory enters generation.

Training uses 3,211 natural Chinese questions with full paragraphs from CMRC 2018 [12]; article overlap is removed across training and development. Both copy and no-copy arms make 4,000 updates from the same V55 weights, retaining natural continuation and motor exercises. On 256 fixed development questions, the copy arm gives 10 exact stopped answers (3.906%), versus one (0.391%) with question pixels blanked and zero for the no-copy control. Raster F1 is 15.718%, 5.553% and 6.582%, respectively (Fig. 4). The paired question-dependent F1 gain is 10.165 points [7.043, 13.371].

Complete reading remains inaccurate. Only one of 393 question pairs sharing a paragraph has both answers correct. Output

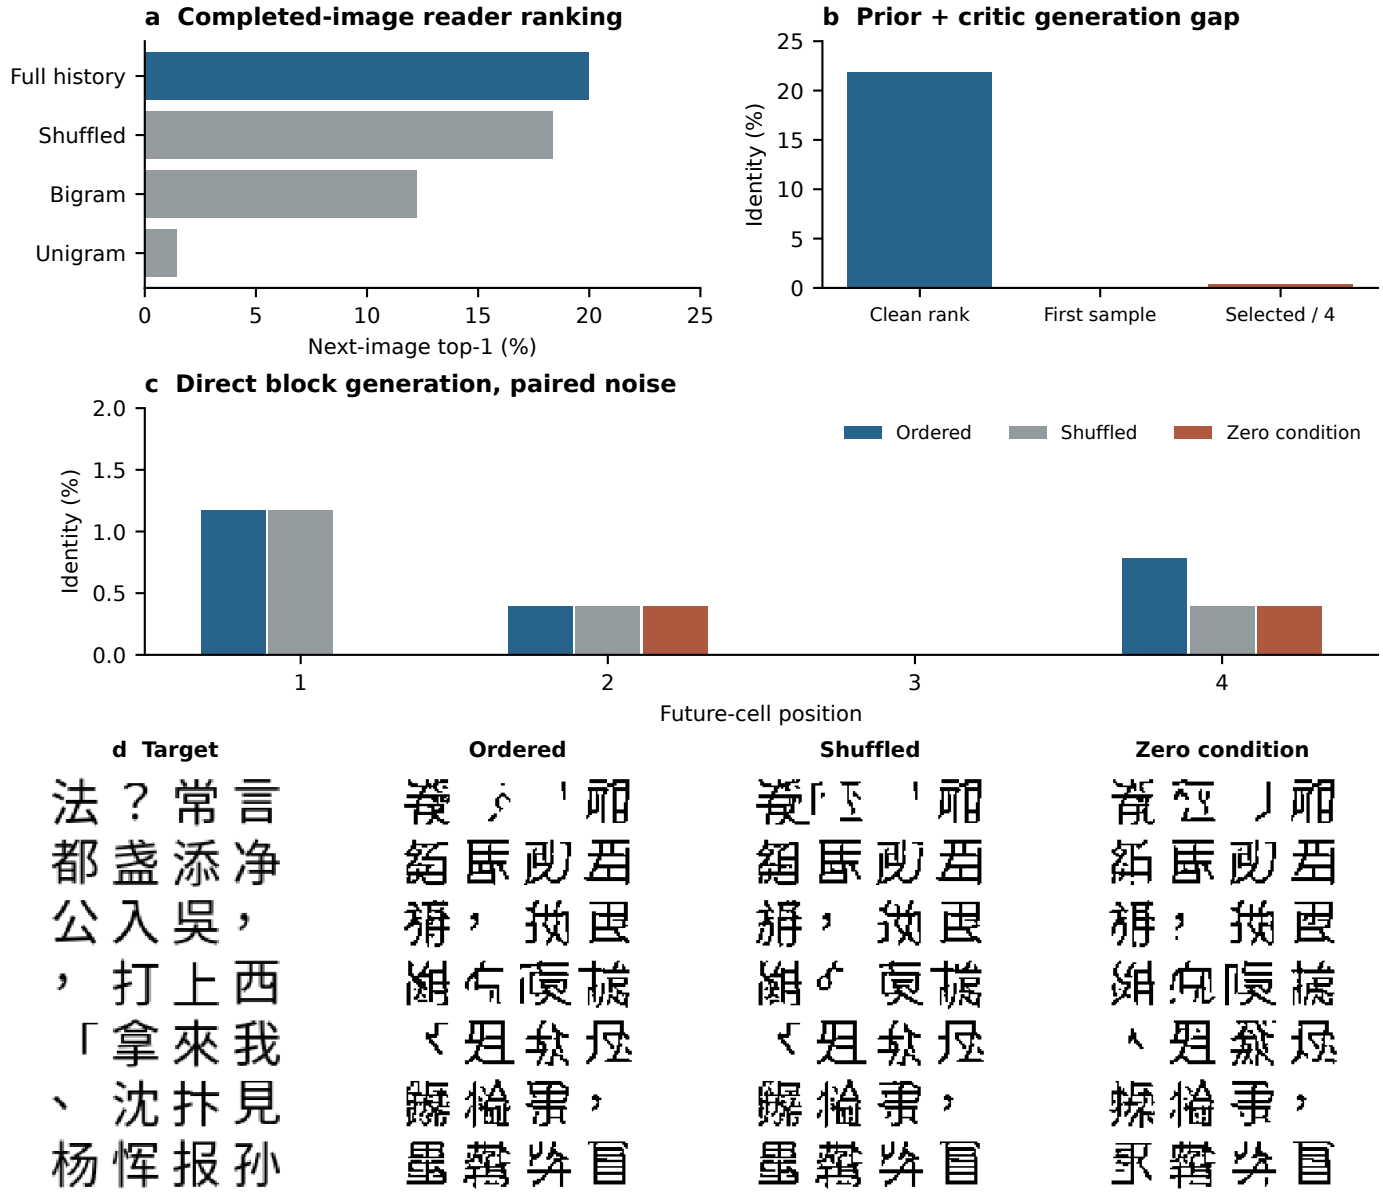


Fig. 2 | Prediction and generated-image evidence. a, The established causal reader ranks completed next images above unigram, bigram and shuffled-history controls on 2,048 windows. These are ranking measurements, not generated text. **b**, Prior-plus-critic guidance retains completed-image ranking but almost never generates the target identity. **c**, Direct conditional block generation reports all four positions under ordered, shuffled and zero conditions. Bars use the same starting noise and single primary samples. **d**, Fixed, unselected examples show the target and all generated controls. Images are raw model outputs with nearest-neighbour enlargement. No OCR or image-bank replacement is applied.

reaches the 17-cell safety cap in 22.656% of primary cases. All 2,677 emitted primary cells use the copy branch; the free writer is retained but not exercised in these answers. All 256 supplied motor images remain exact. Four of seven gates pass, failing answer accuracy, F1 and paired-question correctness; the decision is `reject_reading_model`. A post-result image-only diagnostic finds a correct first glyph in 43 cases, but only 10 become correct complete answers. Fourteen outputs contain a complete reference prefix followed by extra output. Both evidence selection and answer structure need improvement (Supplementary Note 16).

Image-block answers improve structure but retain selection errors

We next normalize a joint start–end law over contiguous observed image blocks of one to sixteen cells. Equal image sequences share probability mass, and the selected end position terminates the answer. Training marginalizes all matching source locations using target images alone. This 19.60M-parameter model retains the free raster writer and natural/motor exercises.

After 4,000 additional updates from the V57 copy checkpoint, exact answers reach 25/256 (9.766%), versus 3/256 with question pixels blanked and 21/256 for matched additional glyph-answer training (Fig. 5). Block F1 is 27.033%, versus 9.169% blanked and 22.140% for the matched control. The block-minus-control F1 gain is 4.893 points [0.469, 9.047], but the exact-answer difference is unresolved: 1.563 [-1.509, 4.652]. Extra training explains much of the improvement over V57; it cannot all be attributed to the block law.

Both answers are correct in only 4/393 paragraph pairs. Six of seven registered gates pass; paired-question correctness fails. Motor copying remains exact on 256/256 images. The block

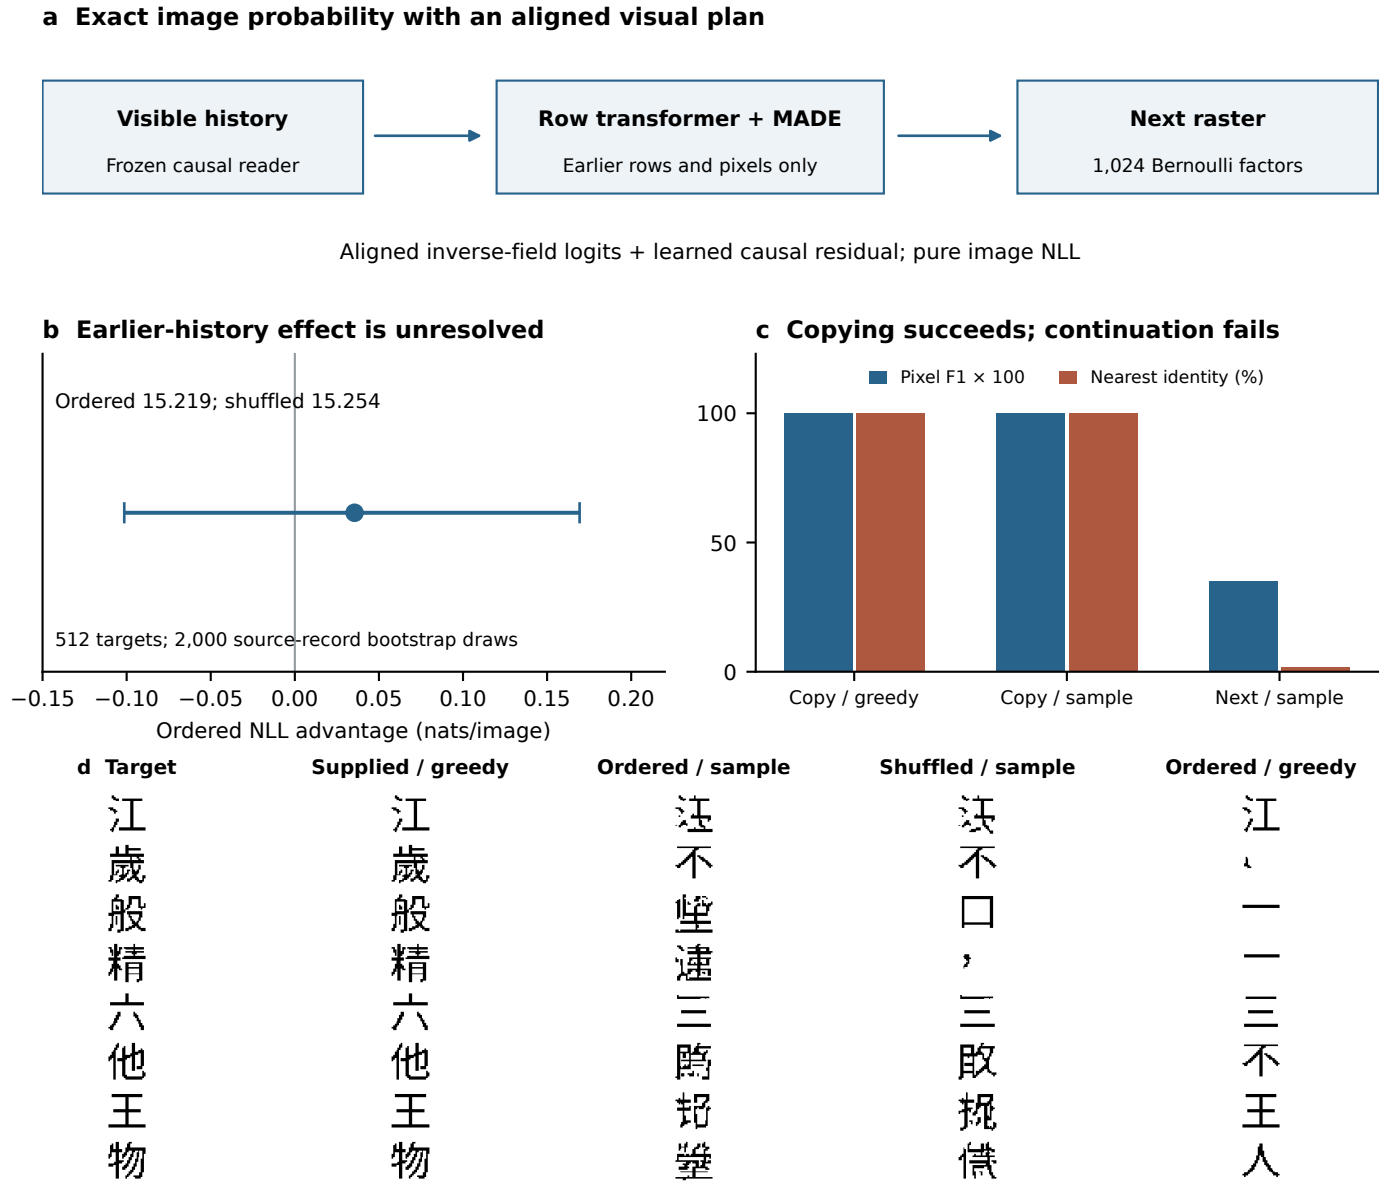


Fig. 3 | Exact image likelihood separates motor success from language failure. **a**, A frozen image reader supplies an aligned plan and continuous history state. Strict row and pixel masks give a normalized raster law. **b**, Ordered-versus-shuffled NLL advantage on 512 development targets; the interval uses 2,000 source-record bootstrap draws and crosses zero. **c**, Supplied-image copying and autonomous next-image sampling on 256 targets. The copy controls reveal the answer. **d**, The first eight audit rows, shown without selection, compare target, exact supplied-image copying and autonomous controls. Every displayed output is the actual model raster. The 9,932-image evaluator bank is created after generation, covers all targets and never selects or replaces an output. Its identity rates are not directly comparable to earlier 1,024-image-bank results.

arm trains in 211.0 s at 0.967 GiB peak allocated memory on the same GPU, excluding previous training. Native prompt-PNG replay through the inference-only export exactly reproduces output pixels and boundary events. The result establishes limited extractive reading, not useful general free writing (Supplementary Note 17).

Frozen pretraining is not sufficient under a shifted visual interface

A separately registered probe compares frozen V58 and released MIXAR-477M[9] features using identical fresh 789,520-parameter block readouts. Both see the same prompt images, pair-aligned states, fixed continuous projections and fine image features. The student readout reaches 24/256 exact answers and 25.340% F1; MIXAR reaches 6/256 and 11.588%. Its F1 difference is −13.752 points [−18.247, −9.094]. No MIXAR pair has both answers correct, versus 12/393 student pairs. The frozen-foundation probe is rejected. This comparison includes different prior task training and a font/layout shift from MIXAR's native renderer; it does not disprove the released model's language competence. External pretraining is not local single-GPU work (Supplementary Note 18 and Fig. S5).

Moving the question before the passage is insufficient

For a causal reader, placing the question first allows each later passage state to depend on it. A registered input-order comparison starts both arms from identical V58 weights and retains the same images, parameters and 4,000-update budget. Question-first reading yields 21/256 exact answers and 23.765% F1, versus 27/256 and 26.908% for continued context-first training. The F1 difference is −3.143 points [−6.544, 0.446], unresolved. Both questions are correct in 5/393 versus 16/393 paragraph pairs.

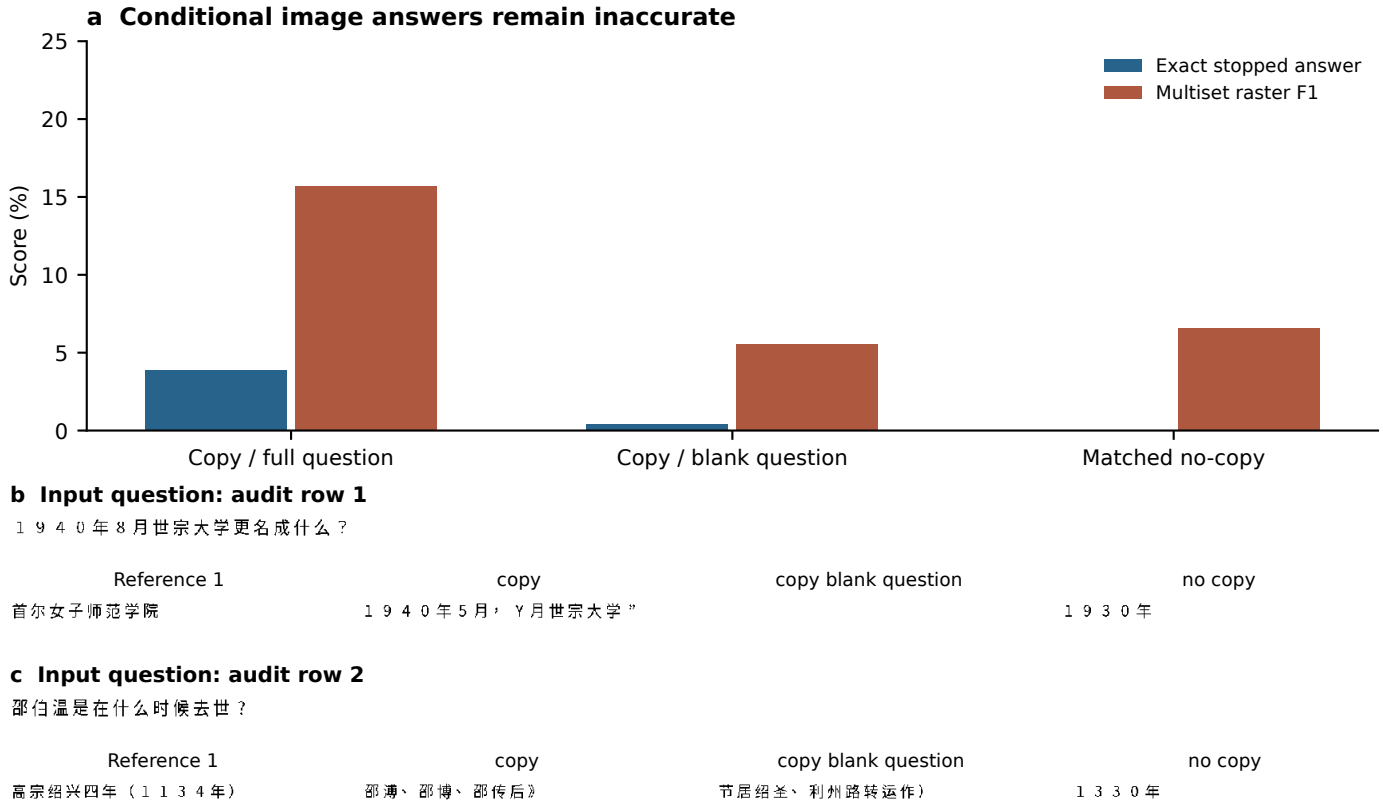


Fig. 4 | Question-dependent image answers remain inaccurate. a, Exact stopped answer and multiset raster F1 on 256 fixed questions from unseen articles, using full questions, blanked question pixels and the matched no-copy arm. F1 does not enforce order or correct termination. **b,c**, First two audit questions and raw answer images, with the first retained reference. The complete paragraph is also supplied to the model. No examples are selected or repaired; white output denotes a zero-cell answer. Question images are host renderings, and outputs are actual model rasters. All primary copy-arm answer cells are selected from the visible prompt.

Four of eight intervention gates pass, and the change is rejected. Motor copying remains exact in both arms. This result limits a simple ordering explanation of the selection failure; it does not establish that question-conditioned source representations are ineffective in general (Supplementary Note 19 and Fig. S6).

Native free writing fits pixels but fails reading

We next adapt every parameter of the released MIXAR-477M native pixel-language model[9], rather than selecting source-image spans. The actor receives prompt images and predicts unrestricted binary image blocks, with no output glyph inventory, OCR or external inference call. All 477,145,984 parameters receive answer-image likelihood gradients. The 1,000-update run takes 1,501.1 s on one RTX 4090 D at 9.912 GiB peak allocated memory, excluding loading and final serialization. External foundation pretraining is attributed separately.

Held-out answer-pixel NLL improves from 0.6842 to 0.1548 nats, yet generated answers remain inaccurate (Fig. 6). Exact answers are 0/256 and both-correct paragraph pairs 0/393. Exact-raster F1 is 0.149%, while aligned foreground F1 is 12.264%. Question blanking reduces foreground F1 to 8.473%, but the strict-raster F1 difference does not resolve zero. Five of ten registered gates pass, all concerning implementation, resources or termination rather than the required answer quality. The decision remains `reject_native_free_reading`.

Native PNG inference reproduces answer tensors, PNG bytes and stopping events from a self-contained learned bundle in a fresh offline process with empty model/code caches. This establishes an independent runtime, without certifying its language. The Noto CJK cell renderer differs from the released model's native font interface. A blinded answer-only OCR diagnostic is poorly calibrated even on reference images, recognizing only 40/256 exactly; it cannot replace a semantic evaluation (Supplementary Note 20).

Correlated pixels improve fit without establishing answers

We next replace the single product-Bernoulli head with eight conditional experts, each defined over a complete 32×64 bitmap. A shared latent expert permits pixel covariance without a character inventory. Both arms fully adapt the same native foundation for 1,000 updates on identical sampled examples; the eight-head actor has 490.00M parameters versus 477.15M for the single-head control. The comparison matches training budget, not capacity.

Held-out answer-pixel NLL is 0.1248 versus 0.1542 nats for eight and one experts, respectively (Fig. 7). Strict-raster F1 increases from 0.234% to 0.815%, a 0.580-point difference [0.089, 1.094]. Foreground F1 increases from 12.205% to 15.827%. Nevertheless, both arms produce zero exact answers among 256 questions and zero both-correct pairs among 393 shared-paragraph pairs. The eight-head normal-minus-blank raster gain is only 0.295 points [−0.006, 0.623], leaving question dependence unresolved. Foreground F1 also has no resolved question-dependent gain. The registered density-intervention gate and both bounded reading gates fail.

Training takes 1,887.5 and 1,972.8 s, with 9.904 and 10.174

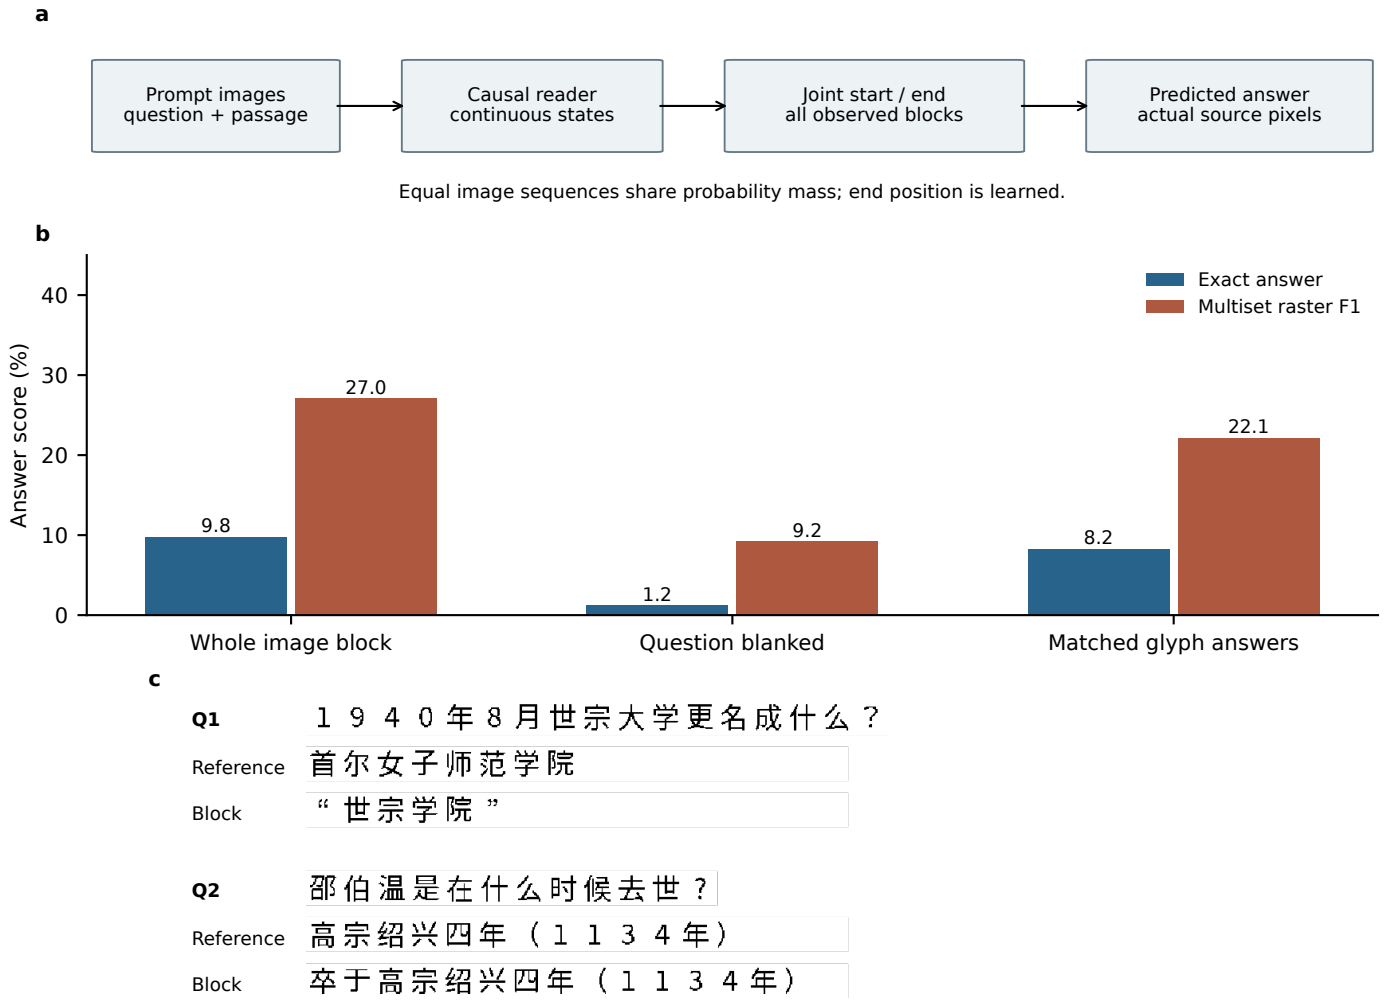


Fig. 5 | Predicting whole observed image answers. a, The reader assigns probability to variable-length source image blocks, marginalizing identical sequences. **b**, Fixed development answers after matched additional training. F1 counts exact-raster multiset overlap; exact answers also require order and termination. **c**, First two audit questions, retained reference images and raw block answers, without example selection or repair. The full passage also enters the model. Block copying preserves forms but does not establish that the selected text answers the question.

GiB peak allocated memory for one and eight experts. Both independent bundles reproduce native PNG tensors, image bytes and decoder events offline. The eight-head gate entropy averages 1.812 nats, with five experts selected across the primary continuations, including termination. These usage diagnostics describe the mixture; they do not demonstrate semantic specialization (Supplementary Note 21).

Direct assignment changes reference scores but weakens free writing

To test conditional binding, we continue the eight-expert generator from one shared learned checkpoint. Two questions about a paragraph provide distinct accepted answer images. Both matched 489,998,208-parameter actors score all four correct and swapped question–answer streams. The control uses positive image likelihood alone; the aligned arm adds a smooth penalty on the assignment margin. Both consume identical pair draws for 2,000 updates. The extra loss cancels additive question-only and answer-only offsets, but this property does not guarantee useful conditional generation.

The aligned actor reaches 2.417% strict-raster F1 versus 3.229% for control (Fig. 8). The difference is −0.812 percentage points, article-bootstrap 95% interval [−1.504, −0.202]. Exact answers are one and two of 256 respectively; both-correct pairs remain zero of 393 for both. Normal-minus-blank-question raster gains resolve positively in each arm: 1.677 points [0.492, 3.033] for aligned and 1.765 [0.504, 3.251] for control. Thus both use question pixels, while neither meets the absolute reading requirements. The assignment intervention and both reading gates fail.

When accepted answer images are supplied, control chooses the correct pair assignment in 324/393 cases and aligned in 330/393. These descriptive scores cannot establish a statistically resolved improvement, and supplied answers are unavailable to autonomous writing. All-development pixel NLL is 0.113901 for control and 0.114186 for aligned. Training takes 7,355.2 and 6,737.3 s, with 10.301 GiB peak allocated memory in both. Complete sample journals match; standalone offline tensors, PNG bytes and stopping events replay exactly (Supplementary Note 22). Timing variation in one sequential run does not establish comparative efficiency.

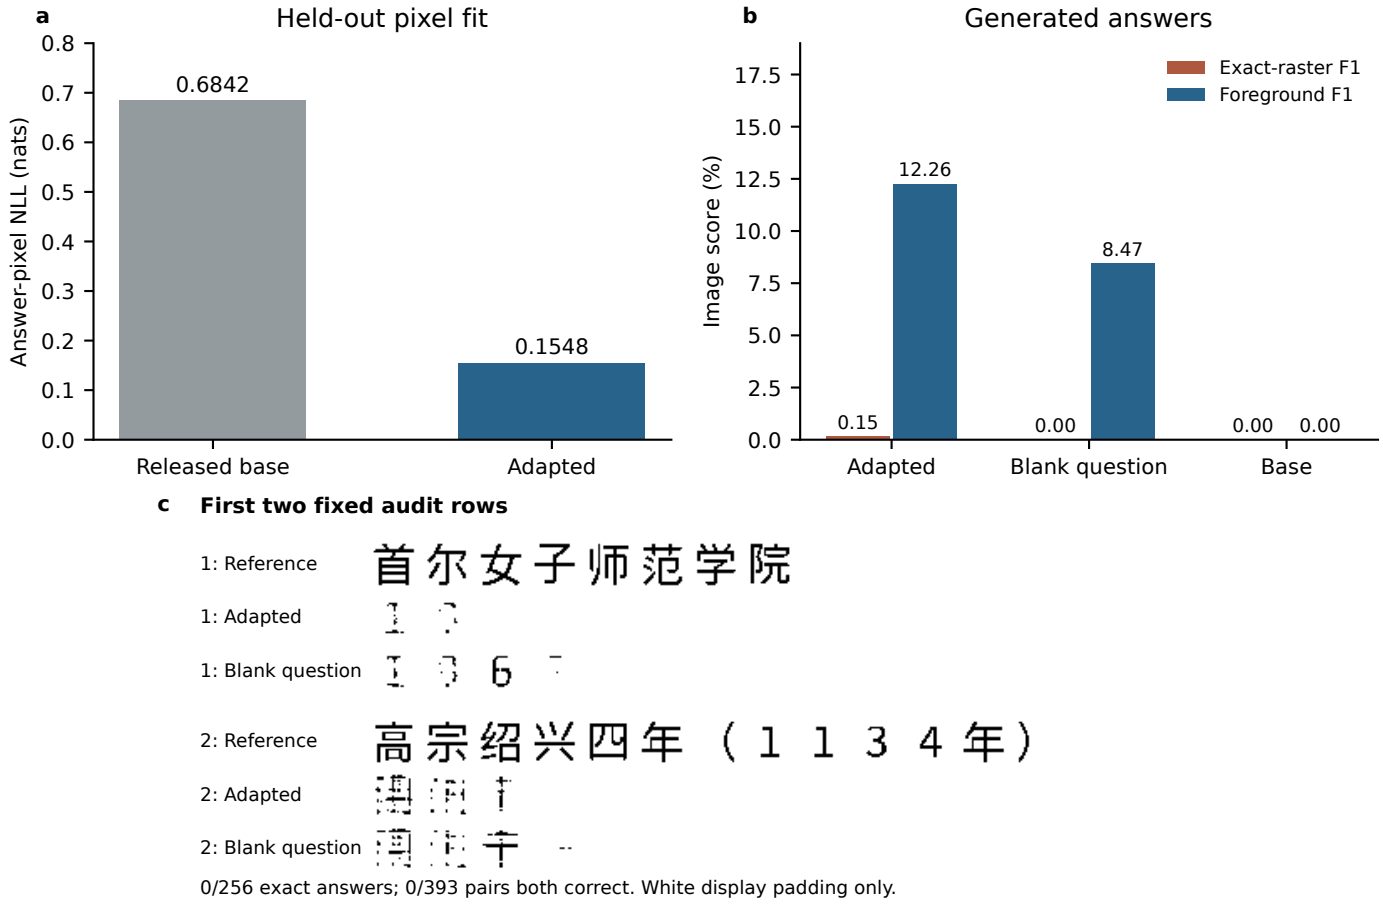


Fig. 6 | Better pixel likelihood does not yield useful free writing. a, All-development answer/termination-pixel NLL, released base and final adapted EMA. **b**, Exact-raster multiset F1 and aligned foreground F1 on 256 fixed primary questions, with question-blanked and released-base controls. Foreground overlap does not establish a correct character or answer. **c**, First two fixed audit rows, with retained references and unmodified generated pixels. White padding aligns the display and never changes a nonblank output. All exact answers and both-correct pairs fail.

Separate output coordinates preserve forms

To isolate a writing failure from conditional language selection, we cross the inherited reader features and a lossless pixel control with flow and direct image reconstruction. All four writers receive the same 64 training blocks, initial weights and 2,000-update exposure. Direct reconstruction recovers 63/64 blocks exactly from inherited features and 64/64 from pixels; neither flow arm recovers an exact block (Fig. 9a,c). This fitting experiment identifies a failure of the tested flow route, but supplies the answer image and cannot demonstrate language.

A broader direct writer then samples 160,000 blocks from a 12,976-example Chinese/English pool, consuming 12,972 distinct examples. After 10,000 updates, the fixed 128-case reconstruction audit reaches 56.563% foreground F1 versus 18.520% with shuffled conditions, with three exact blocks. The conditioning gain passes its original requirement, but fidelity misses 90%. The run uses a fitting-stage warm start and is not a matched broad-corpus comparison with the earlier flow writer. Raw outputs and failed criteria are retained.

A separate continuous writing codec removes the requirement to invert the semantic reader. With its original frozen weights and no additional learning, two 768-dimensional image codes reconstruct the same 128 development blocks at 99.998% F1 and 126 exact blocks, versus 21.413% F1 and zero exact blocks with shuffled codes (Fig. 9b,d). The encoder receives target images; the decoder receives only continuous coordinates. Thus the result qualifies supplied-code fidelity, not prediction of an unknown answer.

The complete successor retains all 27 visual and 36 contextual layers while predicting this separate writing space. It has 7,539,239,282 total parameters and 9,351,937 trainable parameters, with 16.953 GiB peak allocated memory. All 24 actual execution checks pass, including early/late adapter gradients, future-image isolation and exact optimizer/RNG replay. Its two-update output remains fragmented. These preflight measurements establish execution, separately from the subsequent instruction-learning test (Supplementary Notes 24–26).

Conditional prediction remains inaccurate

The complete actor then finishes its registered 10,000-update instruction stage on all 12,976 admitted Chinese/English training examples. Fresh corpus permutations yield 20,000 recorded draws, including every admitted example once in the first epoch. The final adaptation and production checkpoint agree exactly; all recorded draws match the seeded order. This adapts 9,351,937 parameters within the inherited 7.539B-parameter actor, using 16.965 GiB peak allocated memory on one RTX 4090 D. External foundation and codec training remain part of its provenance (Supplementary Note 27).

The final actor generates answers to 128 Chinese questions from 64 shared-passage pairs and 64 English instructions, each with a fixed blank control. All 384 raw streams and input layouts

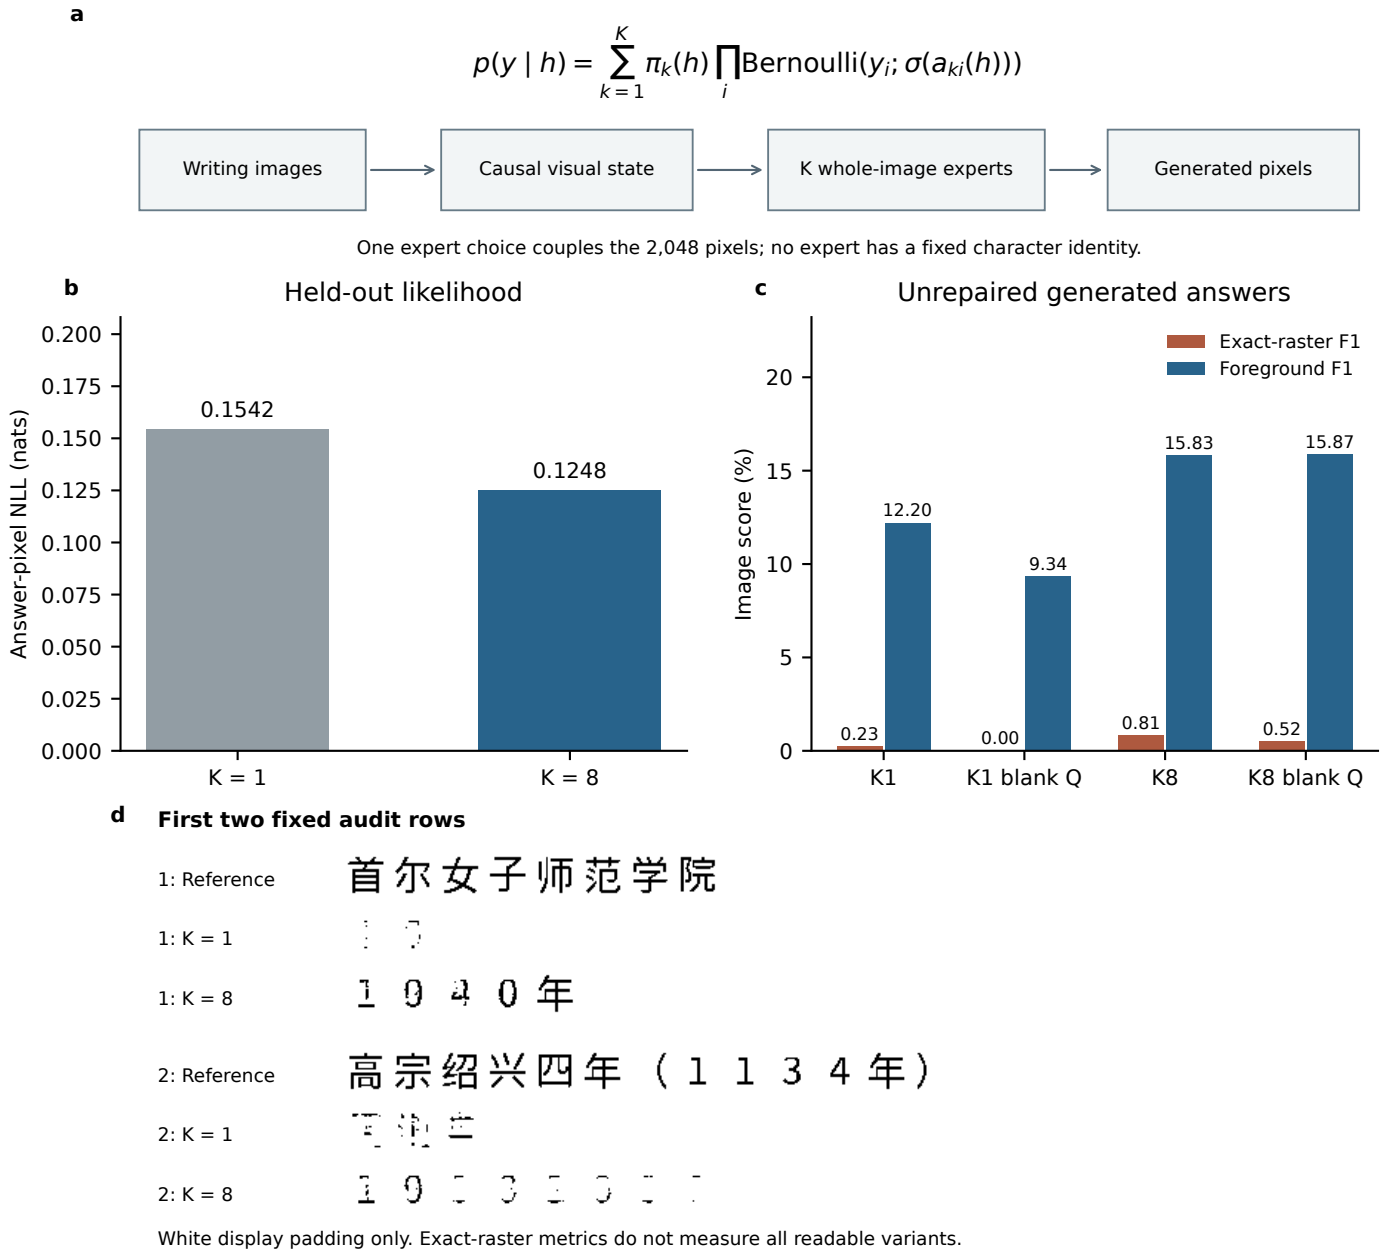


Fig. 7 | A whole-image mixture improves pixel fit but fails the question-dependent writing test. a, Conditional image law and inference boundary. One latent expert couples all pixels; experts are context-dependent neural functions. **b**, All-development answer/termination-pixel NLL after matched-budget adaptation. **c**, Strict-raster and aligned foreground F1 on normal and blank-question prompts. **d**, First two fixed primary references and raw outputs, with white display padding only. No nonblank pixel is repaired or removed. Parameter counts and output computation differ between arms; both use externally pretrained weights.

are independently verified. Chinese foreground F1 is 32.671%, versus 19.508% with question pixels blanked (Fig. 10a). The paired difference is 13.163 percentage points, with 95% article-bootstrap interval [10.933, 15.730]. Nevertheless, no Chinese answer is exact and no pair has both answers correct. All 128 normal Chinese outputs stop through the learned decision. The three absolute answer-quality gates fail while question dependence and stopping pass (Fig. 10b,c). English recorded-reference agreement is also poor, at 9.763% F1 and zero exact matches; that metric does not credit every alternative valid answer.

A fixed 32-case diagnostic then separates supplied-code reconstruction, prediction with exact reference history, and autonomous writing. Each split/source group contains eight examples. Supplied-code F1 is 99.947–100%, but prediction with supplied history and answer length reaches only 21.805–35.554%, including training examples. No first predicted block is exact (Supplementary Note 28 and Supplementary Fig. S9). The initial pixels and stop probabilities agree between paths starting from the same prompt. The failure therefore precedes generated-history drift. This assisted diagnostic does not identify a unique learning repair or qualify language; the original autonomous decision remains rejection.

Numeric controls retain bounded visual reading

We test the original imported body through numeric chat/task vectors with its vocabulary absent. A separate helper temporarily applies the source text projection. Among eight Chinese development questions, native-budget images give four exact answers versus zero with questions erased; original-scale images give one, and matched text gives three. All 742 retained helper

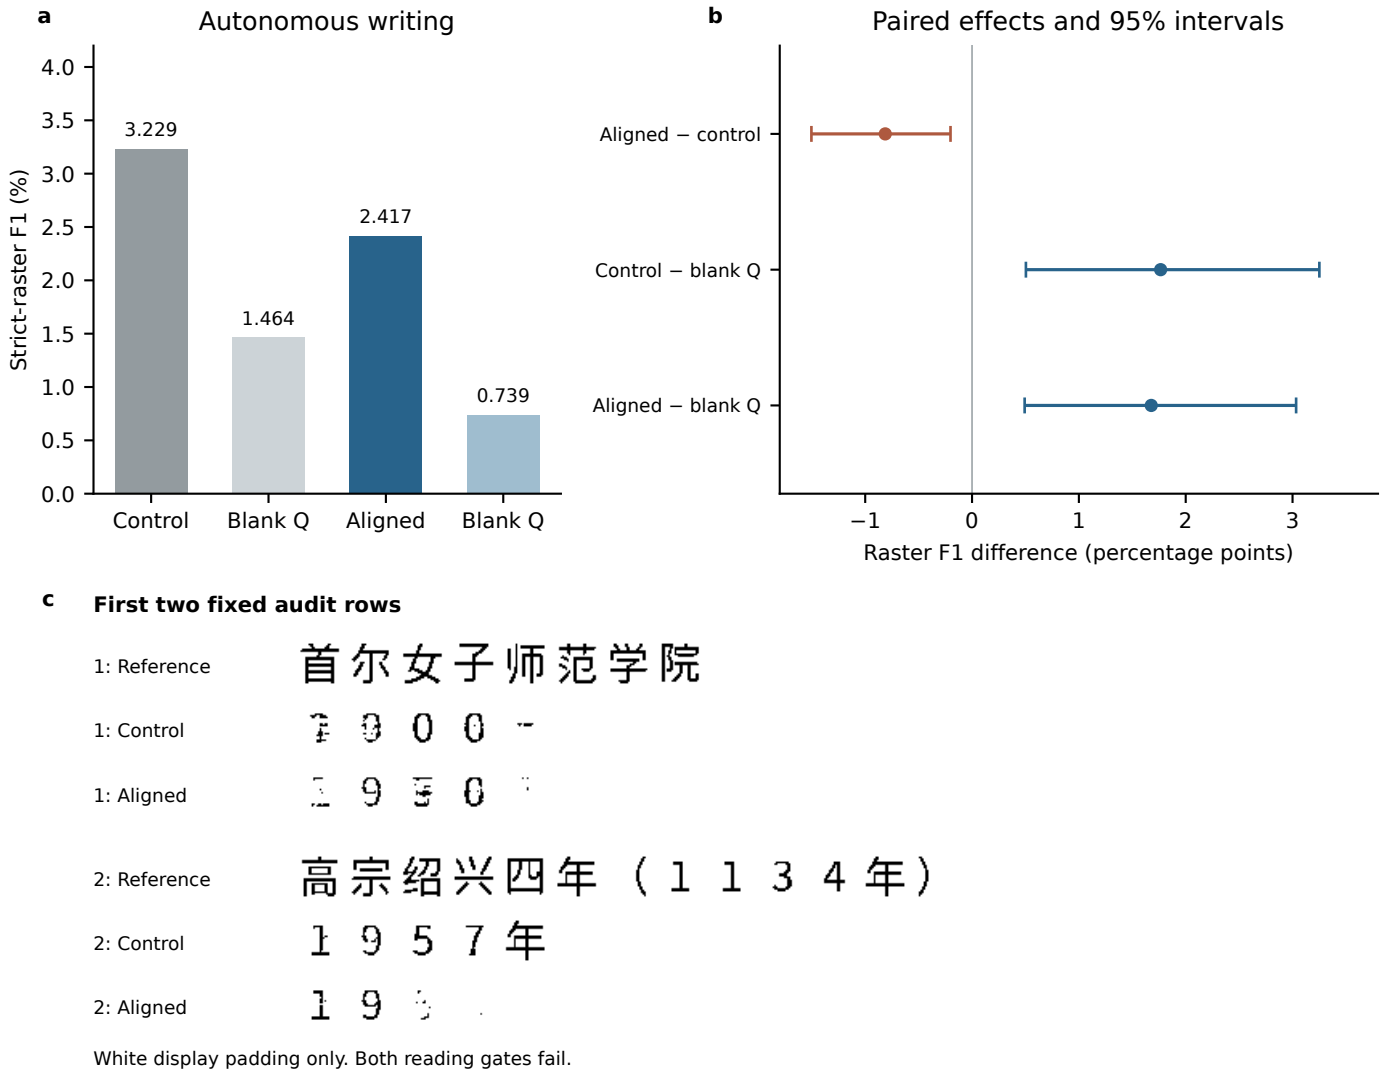


Fig. 8 | Question use and reference compatibility do not establish correct free writing. a, Strict-raster F1 on the fixed 256 primary questions and their blank-question controls. Bars show means. **b**, Paired differences with 95% intervals from 2,000 article resamples. The assignment arm scores lower than control despite positive question effects in both. **c**, First two fixed primary references and unmodified generated answers. White padding aligns the display; no nonblank output pixel is removed or repaired. Strict raster agreement does not credit every readable alternative form.

decisions replay as exact argmax choices (Supplementary Note 29 and Fig. S10). This reused sample supports bounded reading, without qualifying image writing or isolating the scale effect.

The successor conditions a 38.652M-parameter causal writer on frozen prompt-only visual states. A joint flow predicts two-tile codes; the codec emits pixels, which are re-encoded for history. Real-example future-target exclusion, bit-exact optimizer/RNG resume and physical feedback checks pass. Full-corpus training completes 40,000 updates and 640,000 draws over 12,976 examples. The autonomous audit gives zero exact Chinese answers out of 128 and 29.386% foreground F1, versus 21.873% with questions erased. Four of six Chinese requirements fail (Supplementary Note 30 and Fig. S11).

A post-hoc diagnosis identifies an obstruction in the velocity head: its final affine span has rank 769 in 1,536 writing coordinates. In real arithmetic the remaining 767 components of initial noise are invariant before normalization. Actual-weight CPU trajectories support this prediction. A full-dimensional velocity control improves fidelity but fails its worst-image requirement on fresh noise. Predicting clean writing codes instead, with an FP32 final projection and analytic noise contraction, reproduces one learned Chinese image exactly on all 256 fresh draws (Supplementary Note 31). Clean-data prediction follows prior denoising work[19]; the known-image control establishes neither reading nor conditional language competence. The complete clean writer subsequently fits 32 seen training questions: 29 original answer images are exact, versus 31 frozen-codec reference images, with 99.997% original-image F1. This is a paired learning-capacity control, not unseen-question generalization (Supplementary Note 32).

Full-corpus clean writing still fails

A fresh 64.520M-parameter clean writer trains on all 12,976 examples for 40,000 updates, completing 49 corpus passes. Its autonomous audit produces 2/128 exact Chinese answer images and 37.631% foreground F1, versus 23.609% with question ink erased. Both questions are exact in 0/64 passage pairs; three original requirements fail. The question effect is resolved at 14.023 F1 points [10.861, 17.487], and all normal Chinese an-

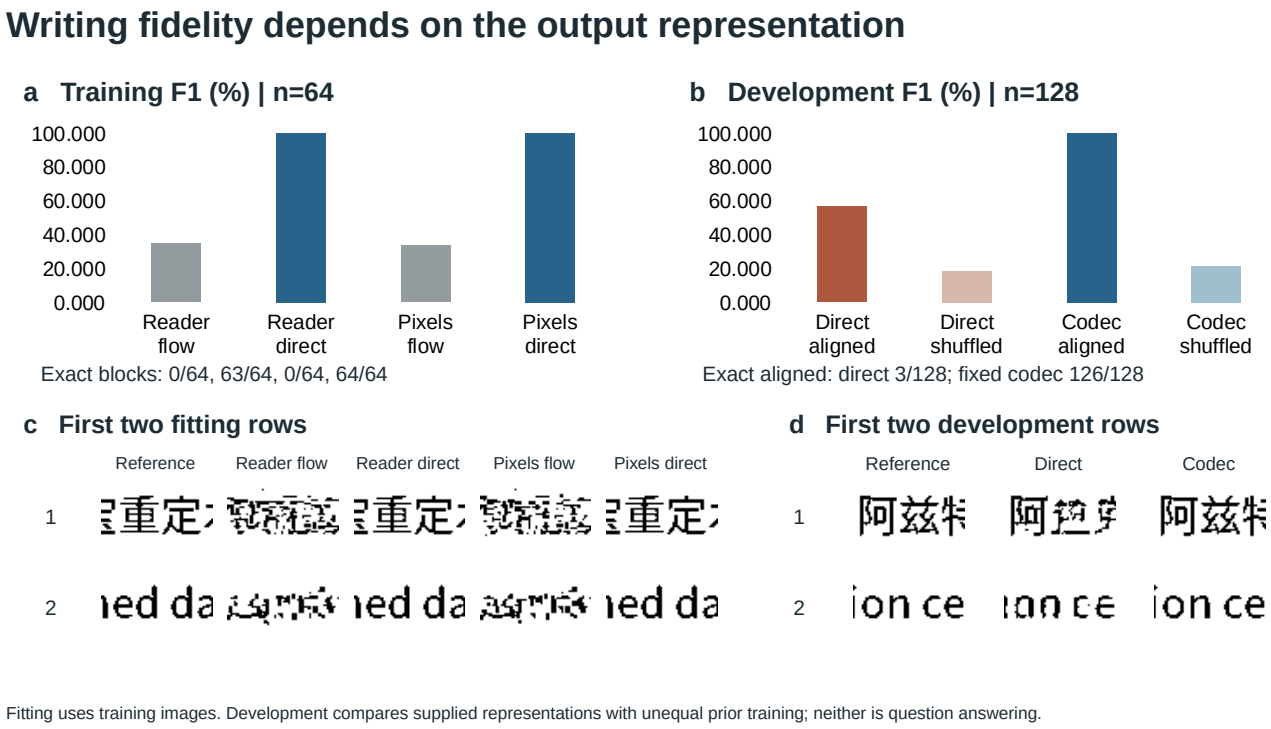


Fig. 9 | Faithful supplied-code writing does not establish answer selection. a, Foreground F1 for four matched fitting arms on 64 training blocks. The pixel condition is an answer-supplying control. **b**, Direct and fixed-codec reconstruction on the same 128 reused-development blocks, with their shuffled-condition controls. Training histories differ; this is not a matched efficiency comparison. **c,d**, The first two fixed fitting and development rows, with original source pixels and unmodified reconstructions. No row is chosen by quality, cropped or repaired.

swers stop by their learned rule. Dependence on question pixels does not establish correct answers. A separate 12-request Chinese/English/numeric copy probe also gives no exact outputs (Supplementary Note 33).

The full writer fails even on the earlier 32 training questions: free generation reaches 45.771% F1 with no exact answer. Correct prior image codes and supplied answer length raise F1 to 58.045%, still with no exact stream. Supplied-code rendering reaches 99.9953% F1. Thus error is already present under correct histories. The clean predictor also slightly corrupts constructed near-clean target codes. This identifies an additional writing defect, without showing that it causes the much larger free-answer failure. The earlier small fit received 2,000 passes over its selected examples; neither that memorization result nor theoretical time-coverage calculations demonstrate a useful general language model.

A matched head-only refinement broadens denoising-time supervision while holding the causal body fixed. It improves assisted near-clean block F1 from 99.0658% to 99.8602%, but free-stream F1 falls from 48.726% to 45.028%. Both arms remain at 0/32 exact training answers. This intervention fails all four prespecified repair criteria (Supplementary Note 34); better near-clean reconstruction is insufficient for prompt-conditioned generation.

An offline source-decoder diagnostic recovers 8/12 exact copy answers and 19/32 exact training answers from the same physical prompt images, compared with zero exact image outputs in both groups (Supplementary Note 35 and Fig. S12). All 44 normal prefill memories match the original 52-state BF16 suffixes exactly. The first source projection matches the correct answer prefix on 22/32 training questions, excluding the paired answer in each case. Thus some answer information is available before the new writer. Later source decoding uses the full layerwise visual/context KV history, which the image writer does not receive. These measurements motivate examining the compact-memory and conditional-readout boundary; they do not demonstrate sufficiency of the compact state or a usable image-language model.

A key-visibility intervention now tests that boundary directly. After identical prefill, permitting attention only to the final 52 task positions reduces exact source-helper copy answers from 8/12 to 0/12 and training answers from 19/32 to 2/32 (Supplementary Note 36 and Fig. S13). Original cache positions, rotary coordinates and first decisions are preserved. This rejects the tested frozen suffix interface without proving a lower bound on learned compression. The next continuous image actor retains all 36 native layers and full prompt KV. Initial forward, gradient, causal-mask and emitted-image feedback checks use 16.638 GiB on one GPU, with zero optimization updates. Its first writing block still depends on the initial reader state alone. The complete trained actor has 7.583B parameters including externally pretrained weights; 46.900M image interface parameters are trainable.

Subsequent training completes 40,000 updates over all 3,193 Chinese examples, using 18.298 GiB peak allocated memory, but gives 0/128 exact development images, 36.504% foreground F1 and 0/64 both-correct pairs (Supplementary Note 37 and Fig. S14). Erasing the question reduces F1 to 25.404%; the paired gain is 11.101 percentage points, with a 95% article-bootstrap interval of [8.578, 14.018]. All streams stop, but stopping and input sensitivity do not establish accurate answers. On the same 32 training questions, free generation gives 46.235% F1 and correct-history prediction gives 50.241%, with no exact

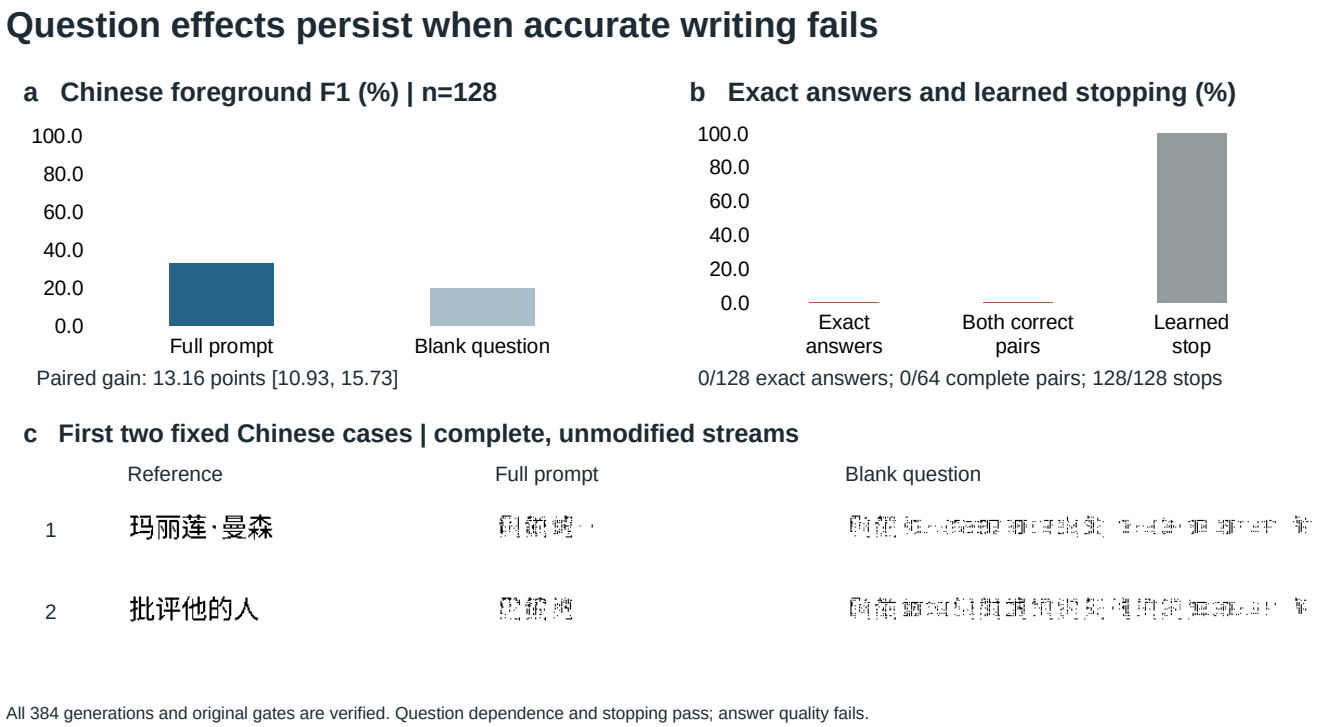


Fig. 10 | Question dependence persists when accurate answer writing fails. a, Chinese foreground F1 on 128 fixed development questions with full and matched question-blank prompts. The paired difference uses 2,000 bootstrap draws of 64 source articles. **b**, Exact full answers, both-correct question pairs and learned stopping. The original answer-quality gates fail, although question dependence and stopping pass. **c**, Complete reference and unmodified generated streams for the first two fixed Chinese cases. Images retain their physical width and aspect ratio; no nonblank pixel is removed or repaired. These are reused development data.

images in either path. Supplied-code decoding gives 99.9953% F1 and 28/32 exact images. Thus conditional prediction remains inaccurate under correct history; generated-history errors cannot be the sole explanation. These physical-image scores are not a separate measurement of semantic accuracy.

A matched readout repair adds an initial continuous query and trains rank-8 suffix adapters while retaining the full prompt cache. After 2,000 updates per arm on those same 32 training questions, free F1 is 71.177% for continued control and 72.748% for adaptation; neither produces an exact complete image (Supplementary Note 38 and Fig. S15). Correct preceding history gives 77.372% and 79.629%, again with no exact images. Supplied-code fidelity is unchanged at 99.9953%. The adapted model fails three of four declared training-fit criteria. The small increase in image agreement does not resolve conditional writing or qualify general language capability.

Decoder-sensitive reweighting also fails (Supplementary Note 39 and Fig. S16). From identical adapted weights, 2,000 matched updates lower autonomous F1 from 85.984% to 70.009%, with 0/32 exact in both arms. Correct-history F1 is 92.430/75.116%; supplied codes give 99.9953%. All four repair criteria fail; scaling is rejected.

Discussion

The results separate a useful visual predictive state from a useful visual language generator. The established causal reader learns context-dependent next-writing information, while subsequent writers can erase that information or model the marginal appearance of glyphs. A complete visual language model must bind meaning and sequence to the exact marks it generates.

An image interface enlarges the representable output space, including the possibility of forms absent from a character encoding. It does not remove the need to learn syntax, knowledge or long-range dependencies. This study has not demonstrated general instruction following, etymology answers, handwriting transfer or arbitrary historical-script generation. It also lacks independent training-seed replication and a capability-matched text-model comparison. Small allocated memory in a bounded experiment establishes feasibility on one GPU, rather than a general efficiency advantage over language models.

Supplied-target reconstruction is preserved by an aligned inverse-image path; new prompt-directed copying still fails. The remaining test is whether context selects and generates a coherent written future. Low teacher-forced image NLL can reflect dependencies among strokes already revealed within the target, while generated output still loses global identity. Joint training improves proper likelihood beyond matched extra density training; the raw generated continuations still fail the intended language test. Conditional continuation and short answers require counterfactual prompts and the same context effect during autonomous sampling and visible feedback. Broader books and offline teacher-generated exercises can supply later training data while leaving the final inference path independent. A larger corpus or longer training requires improvement on the intended generation task, beyond the conditional-likelihood gain established here.

Recent visual generators study mixture assignment and the geometry of distribution matching [13,14]. These results motivate more expressive writing distributions, but marginal image fidelity cannot certify an answer. A generator that swaps the correct

images for two equally frequent questions preserves the answer-image distribution while answering both incorrectly. Conditional supervision and prompt interventions therefore remain necessary; natural-image gains do not establish a shortcut to language learning.

An inherited visual/contextual successor provides a separately measured execution route. With its vocabulary module removed, the complete 7.548B actor trains 25.58M adapter/writer parameters within 17.181 GiB allocated memory. Causal, backward and resume checks pass after explicit precision repairs. Eight short source-engine sanity cases support the contextual initialization; no image-language ability follows from that check. Its continuous writer fails its reconstruction test (Supplementary Note 23). The separate qualified codec preserves supplied writing, while the completed predicted-code actor fails its autonomous language test (Supplementary Notes 25–28). Poor first-block and exact-history predictions, including on training cases, show that generated-history drift cannot be the sole explanation. Conditional feature information, readout capacity, the objective and optimization remain possible limitations. The current evidence does not select a successful density, optimizer or semantic-transfer method. External pretraining remains part of the model's provenance.

Methods

Data, rendering and partitioning

The public Chinese stream contains 7,017 provenance-tracked records. Records are assigned deterministically to training, development or a sealed partition before rendering. Original and simplified views inherit their source split. Both generation experiments use one fixed Noto Sans CJK font, 26-pixel writing inside 32-pixel cells, without image augmentation. A 65-cell training segment supplies the first 60 cells as context and cells 61–64 as future writing. The remaining cell is unused. No string, source identifier or symbolic target enters a model method. Sealed records remain unopened.

Training and numerical verification

In the diffusion experiment, the reader is frozen. Copied glyph-prior weights initialize the denoiser, which is then optimized as a velocity predictor. AdamW uses a learning rate of 5×10^{-5} for transferred layers and 2×10^{-4} for new modules, 400 warmup updates, cosine decay to 10%, batch size eight, BF16 arithmetic and gradient clipping at one. EMA decay is 0.999. Checkpoints retain model, optimizer, RNG, dataset offset and provenance hashes. Interrupted and continuous smoke runs produce byte-identical model and EMA weights.

The field objective applies an orthonormal transform directly to floating signed predictions. Thresholding is confined to output and identity evaluation. This distinction is verified by an isolated pixel-gradient test. The primary velocity MSE is supplemented by foreground, Dice, field and representation losses at weight 0.10 each, plus an ordered-history margin at weight 0.25. Supplementary Information gives the complete objective and gates.

The exact-density experiment uses 64-cell causal histories and 16 stratified positions in each of eight streams per update. A quarter of selected images receive supplied-image motor conditions; the remainder use predicted states. Binary observations are thresholded once before unweighted Bernoulli NLL. AdamW uses learning rate 2×10^{-4}, 400-update warmup, cosine decay, BF16, norm clip one and EMA 0.999. All 8,000 updates complete in 207.8 s with 0.400 GiB peak allocated CUDA memory. This timing excludes the frozen reader's earlier training. Exact resumed weights and EMA match uninterrupted smoke training. Strict-causality gradient tests and enumeration of a four-bit subproblem verify the probability masks.

The matched continuation reinitializes optimizer and EMA from identical exact-density weights. Both arms make 8,000 additional updates, using the same data and training generator. The joint arm uses reader learning rate 2×10^{-5} and density rate 2×10^{-4}, with the same schedule and motor mixture. It trains 18,711,840 participating parameters; the frozen arm trains 2,434,592. Joint and control times are 342.4 and 218.0 s, with peaks 0.648 and 0.400 GiB. Frozen-reader weights remain identical to initialization, while joint reader weights change. Both final training RNG states match. Development opens only after both arms finish.

The subsequent decoder retains eight partial images after each binary pixel expansion, ranked by accumulated Bernoulli log probability. Row conditions are gathered along beam-parent indices within each row. Stable ties prefer ink, matching the original width-one rule. Full FP32 matmuls are used for both new controls after a documented numerical preflight; model weights are unchanged. The maximum production incremental-score/parallel-NLL error is 2.10×10^{-5} nat, below the unchanged 0.001 tolerance. The complete inference audit takes 269.3 s at 0.154 GiB allocated CUDA memory.

The assignment experiment keeps the native eight-expert architecture and trains all parameters from the same V62 EMA. Both arms score the same four streams for every sampled pair. AdamW, sample order, update budget and final EMA selection match. The auxiliary penalty has temperature 0.01 and weight one for aligned, zero for control. Complete details, independent runtime replay and the unchanged conjunctive gates appear in Supplementary Note 22.

Separate continuous writing coordinates

The successor uses a visual reader E_r, contextual body F_θ, output projection A_θ and frozen writing codec (E_w, D_w). For observed prompt image X and completed answer blocks $Y_{<t}$,

$$h_t = F_\theta(E_r(X), E_r(Y_{<t})), \tag{5}$$

$$\widehat{z}_t = N(A_\theta h_t), \qquad z_t^* = \mathrm{sg}(E_w(Y_t)), \tag{6}$$

$$\widehat{Y}_t = D_w(\widehat{z}_t). \tag{7}$$

Here h_t contains two structural queries, N is nonaffine LayerNorm with $\varepsilon = 10^{-5}$, and sg stops label gradients. Each writing state comprises two continuous 768-vectors for two physical 32×32 tiles, without character segmentation. The encoder already normalizes its codes; the state loss explicitly normalizes both arguments again:

$$\begin{aligned} \mathscr{L} = {} & \mathrm{MSE}(N(\widehat{z}_t), N(z_t^*)) \\ & + \mathrm{MSE}(D_w(\widehat{z}_t), Y_t) + \mathscr{L}_{\mathrm{stop}}. \end{aligned} \tag{8}$$

Uniformly sampled block losses retain the complete teacher-forced history. A separate class-balanced stop loss supervises termination after the complete answer. Frozen decoder weights still transmit image-loss gradients into predicted codes and contextual adapters. Generation uses prompt pixels, learned stopping and only its own completed output images. This regression objective is not a normalized language likelihood or a guarantee

against averaging distinct conditional writing modes. Execution and reconstruction are evaluated separately from language (Supplementary Notes 26–28). The complete instruction run fixes 10,000 updates and two examples per update under fresh full-corpus permutations, with final raw weights and no EMA. Its autonomous audit scores all 192 fixed normal prompts and their matched blank controls, preserving full widths and learned stops. Neither the actor nor its evaluator repairs generated pixels. The subsequent exact-history diagnostic uses the same final weights and no optimizer updates.

Sampling, controls and statistics

The diffusion experiment uses 32 registered DDIM schedule levels and classifier-free scale two. Four future cells start from Gaussian noise. The same noise is reused for ordered, suffix-preserving shuffled and zero-condition histories. Only one sample contributes to the primary block result. Generated cells are appended as visible rasters for a second four-cell block.

An evaluator-only canonical bank measures nearest-raster identity after generation. Targets outside the bank receive no identity credit, and bank coverage is reported per position. This measurement can assign a nearest identity to an illegible pseudo-glyph, so direct images and surface statistics accompany every result. Pixel F1 alone is also insufficient because many Chinese glyphs share strokes. No OCR is used by the model or the reported identity evaluator.

The exact-density audit uses 512 NLL targets, 256 single-image generations and 64 four-step feedback streams. Its sampler uses 32 row-transformer evaluations and 1,024 masked-head evaluations per image. Ordered and shuffled runs share the same Bernoulli uniforms. Its post-generation bank includes all distinct training forms plus every audit target, including punctuation: 9,932 unique binary rasters and 100% target coverage. This enlarged bank differs from the earlier Han-only bank. Zero condition is an out-of-distribution intervention, not a separately trained unconditional baseline. One standalone export embeds the complete reader and density; image-prompt inference reproduces the training-checkpoint continuation exactly under the same seed and greedy mode.

Confidence intervals use 2,000 bootstrap resamples of source records, retaining all windows from a record together. Paired differences use identical windows and initial noise. Intervals describe development-record variability for this trained checkpoint, not variability across training seeds. The current development partition has informed a sequence of experiments and is not a final independent generalization test.

The assignment comparison retains the same fixed 256 primary questions and 393 paragraph pairs. Its paired 95% intervals use 2,000 whole-article resamples, seed 20266329. Reference-energy ranking supplies accepted answers and remains a separate diagnostic. One seed, reused development data and external foundation training limit both this result and efficiency claims.

Data and code availability

Implementation, protocols, evidence summaries and scripts are available in `https://github.com/lachlanchen/ImagizedLanguageModel`. Manifest and checkpoint hashes identify the exact local research artifacts. Restricted reference books are not redistributed. The complete research record and Supplementary Information distinguish earlier foundation provenance from the present raster-only deployment experiment.

References

1. Tai, Y., Liao, X., Suglia, A. & Vergari, A. PIXAR: Auto-Regressive Language Modeling in Pixel Space. *arXiv* 2401.03321 (2024). `https://arxiv.org/abs/2401.03321`.
2. Li, T., Tian, Y., Li, H., Deng, M. & He, K. Autoregressive Image Generation without Vector Quantization. *Advances in Neural Information Processing Systems* (2024). `https://arxiv.org/abs/2406.11838`.
3. Chen, B. et al. Diffusion Forcing: Next-token Prediction Meets Full-Sequence Diffusion. *arXiv* 2407.01392 (2024). `https://arxiv.org/abs/2407.01392`.
4. Dhariwal, P. & Nichol, A. Diffusion Models Beat GANs on Image Synthesis. *Advances in Neural Information Processing Systems* (2021). `https://arxiv.org/abs/2105.05233`.
5. Yu, S. et al. Representation Alignment for Generation: Training Diffusion Transformers Is Easier Than You Think. *International Conference on Learning Representations* (2025). `https://arxiv.org/abs/2410.06940`.
6. van den Oord, A., Kalchbrenner, N. & Kavukcuoglu, K. Pixel Recurrent Neural Networks. *Proceedings of the 33rd International Conference on Machine Learning* **48**, 1747–1756 (2016). `https://proceedings.mlr.press/v48/oord16.html`.
7. van den Oord, A. et al. Conditional Image Generation with PixelCNN Decoders. *arXiv* 1606.05328 (2016). `https://arxiv.org/abs/1606.05328`.
8. Germain, M., Gregor, K., Murray, I. & Larochelle, H. MADE: Masked Autoencoder for Distribution Estimation. *Proceedings of the 32nd International Conference on Machine Learning* **37**, 881–889 (2015). `https://arxiv.org/abs/1502.03509`.
9. Hu, C., Tai, Y., Vergari, A., Keller, F. & Suglia, A. MIXAR: Scaling Autoregressive Pixel-based Language Models to Multiple Languages and Scripts. *arXiv* 2604.11575v2 (2026). `https://arxiv.org/abs/2604.11575v2`.
10. Holtzman, A., Buys, J., Du, L., Forbes, M. & Choi, Y. The Curious Case of Neural Text Degeneration. *International Conference on Learning Representations* (2020). `https://arxiv.org/abs/1904.09751`.
11. See, A., Liu, P. J. & Manning, C. D. Get To The Point: Summarization with Pointer-Generator Networks. *Proceedings of ACL*, 1073–1083 (2017). `https://aclanthology.org/P17-1099/`.
12. Cui, Y. et al. A Span-Extraction Dataset for Chinese Machine Reading Comprehension. *Proceedings of EMNLP-IJCNLP*, 5883–5889 (2019). `https://aclanthology.org/D19-1600/`.
13. Zhang, C. et al. Unifying Distributional Training for One-Step Visual Generation. *arXiv* 2609.35763v4 (2026). `https://arxiv.org/abs/2609.35763v4`.
14. Van Assel, H., De Brouwer, E., Saremi, S., Scalia, G. & Regev, A. Generate in Reconstruction Space, Match in Semantic Space: Transport Geometry for One-Step Generation. *arXiv* 2606.00514v2 (2026). `https://arxiv.org/abs/2606.00514v2`.
15. Tao, H. et al. Multimodal Flow: Unified Flow Modeling of Language and Vision in Embedding Spaces. *arXiv* 2609.40362v1 (2026). `https://arxiv.org/abs/2609.40362v1`.
16. Huang, L. & Wu, X. OCG: Optical Character Grounded Visual Reasoning with Flow Matching. Author manuscript (2026). Peer-review status unverified.
17. Bai, S. et al. Qwen3-VL Technical Report. *arXiv* 2511.21631v2 (2025). `https://arxiv.org/abs/2511.21631v2`.
18. Databricks. Databricks Dolly 15k instruction dataset (2023). `https://huggingface.co/datasets/databricks/databricks-dolly-15k`.
19. Li, T. & He, K. Back to Basics: Let Denoising Generative Models Denoise. *arXiv* 2511.13720v1 (2025). `https://arxiv.org/abs/2511.13720v1`.